

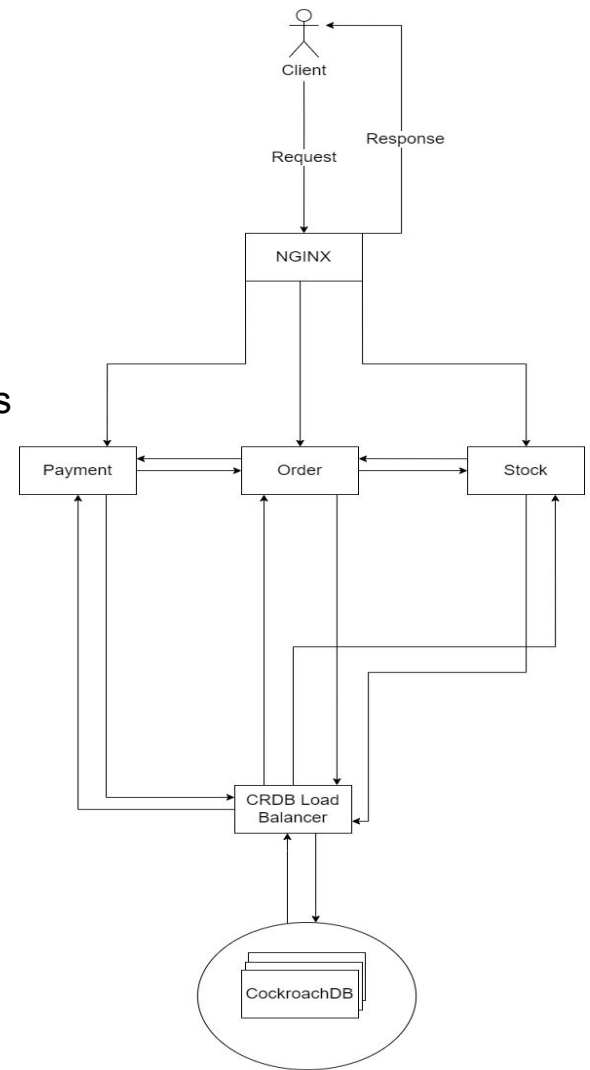
Developing Scalable Services and Managing their Transactions with CockroachDB



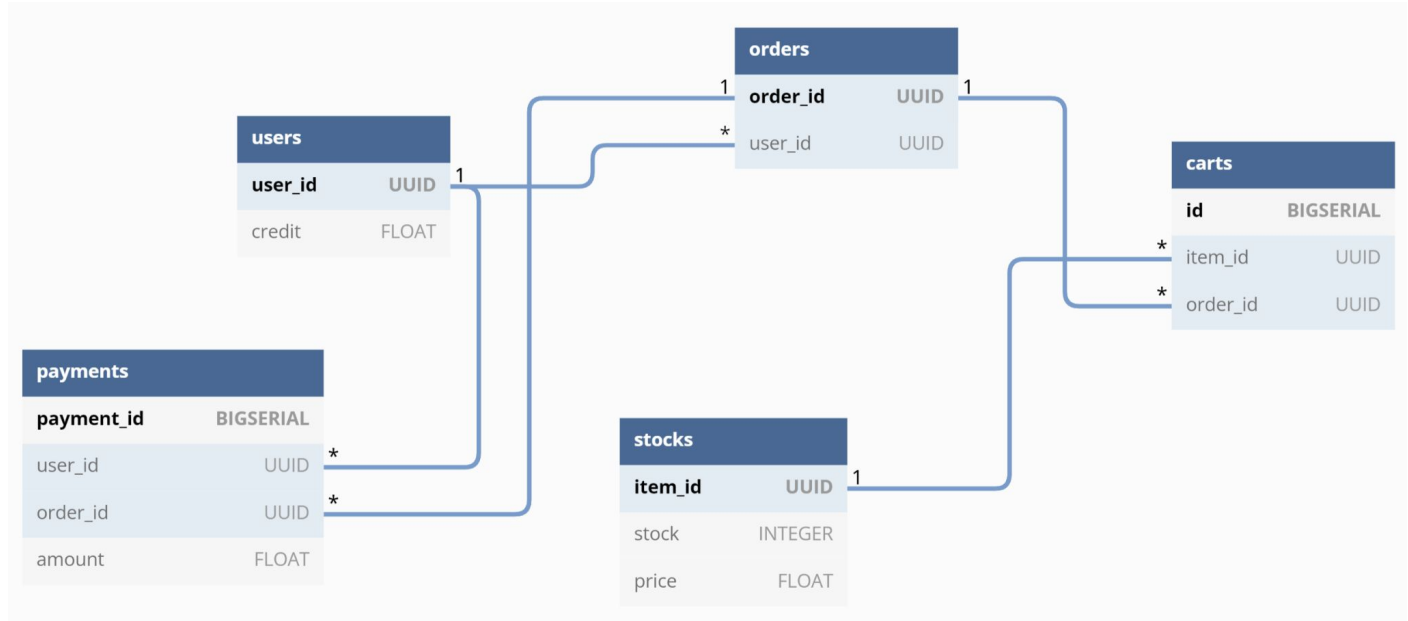
Adam Kadiev
Ayush Patandin
Jakub Nguyen
Mohamed Rashad
Kevin Nanhekhan

System Design

- Services are stateless
 - Each service can scale independently
- Interservice communication through REST calls
 - Services use Flask API
- Database used is a relational database
 - Key-value store under the hood
 - Distributed database
 - Load balancer
 - Adheres to ACID principles



CockroachDB table overview




Method of transaction execution

- SQLAlchemy ORMs
 - SQLAlchemy sessions
 - Model Schema (Python Classes)
- Transactions using Two-Phase Commit
 - Services become stateful
 - Difficult to scale
- Transactions Managed by database
 - Services are stateless
 - Easy to scale

Method used for consistency

- Two-Phase Commit
 - Slow due to synchronous REST calls
- Transactions
 - Automatic retries
 - ACID guarantees



Scalability

- Horizontal Pod Scaling in Kubernetes
 - Given CPU usage, creates new instance of services
 - Minimum and maximum amount of pods specified

Fault-tolerance

- CockroachDB replication and partitioning

Results: 2PC

- Delayed requests due to DB block
- 0 failures for 5 spawned users
- Prepared transactions in checkout




| Type | Name | # Requests | # Fails | Median (ms) | 90%ile (ms) | 99%ile (ms) | Average (ms) | Min (ms) | Max (ms) | Average size (bytes) | Current RPS | Current Failures/s |
|------------|--|------------|---------|-------------|-------------|-------------|--------------|----------|----------|----------------------|-------------|--------------------|
| POST | /orders/addItem[{order_id}/{item_id}] | 71 | 0 | 17 | 25 | 37 | 18 | 13 | 37 | 0 | 0.3 | 0 |
| POST | /orders/checkout[{order_id}] | 35 | 0 | 100 | 180 | 250 | 120 | 78 | 249 | 7 | 0.3 | 0 |
| POST | /orders/create[{user_id}] | 37 | 0 | 18 | 27 | 62 | 21 | 13 | 62 | 52 | 0 | 0 |
| POST | /payment/add_funds[{user_id}/{amount}] | 37 | 0 | 20 | 32 | 48 | 22 | 13 | 48 | 14 | 0.1 | 0 |
| POST | /payment/create_user | 39 | 0 | 16 | 21 | 32 | 17 | 13 | 32 | 51 | 0.2 | 0 |
| POST | /stock/add[{item_id}/{number}] | 80 | 0 | 17 | 27 | 47 | 20 | 14 | 47 | 0 | 0.7 | 0 |
| POST | /stock/item/create[{price}] | 80 | 0 | 15 | 28 | 100 | 20 | 12 | 102 | 51 | 0.7 | 0 |
| Aggregated | | 379 | 0 | 17 | 62 | 180 | 29 | 12 | 249 | 23 | 2.3 | 0 |

Results: Scalable Version

- 4 order and payment deployments
- 1 stock deployment
- Faster than 2PC
- Few inconsistencies for ~1000 users
→ resolvable by performing single transaction in checkout

```
Consistency test - Creating tmp folder...
Consistency test - tmp folder created
Consistency test - Populating the databases...
populate - Creating items ...
populate - Items created
populate - Creating users ...
populate - Users created
Consistency test - Databases populated
Consistency test - Starting the load test...
stress - Creating orders...
stress - Orders created ...
stress - Running concurrent checkouts...
stress - Concurrent checkouts finished...
Consistency test - Load test completed
Consistency test - Starting the consistency evaluation...
verify - Stock service inconsistencies in the logs: -4
verify - Stock service inconsistencies in the database: 0
verify - Payment service inconsistencies in the logs: 4
verify - Payment service inconsistencies in the database: 4.0
Consistency test - Consistency evaluation completed
```



| Type | Name | # Requests | # Fails | Median (ms) | 90%ile (ms) | 99%ile (ms) | Average (ms) | Min (ms) | Max (ms) | Average size (bytes) | Current RPS | Current Failures/s |
|------------|--|------------|---------|-------------|-------------|-------------|--------------|----------|----------|----------------------|-------------|--------------------|
| POST | /orders/addItem{order_id}/{item_id} | 2312 | 0 | 360 | 1800 | 3900 | 678 | 17 | 6562 | 0 | 7.7 | 0 |
| POST | /orders/checkout{order_id} | 1497 | 0 | 700 | 2100 | 4800 | 958 | 44 | 6728 | 10 | 6 | 0 |
| POST | /orders/create{user_id} | 1530 | 0 | 340 | 1700 | 3600 | 657 | 16 | 6538 | 52 | 4.7 | 0 |
| DELETE | /orders/removeItem{order_id}/{item_id} | 325 | 0 | 420 | 2000 | 3900 | 761 | 18 | 6492 | 0 | 1.5 | 0 |
| POST | /payment/add_funds{user_id}/{amount} | 1325 | 0 | 73 | 190 | 440 | 94 | 18 | 650 | 14 | 3.4 | 0 |
| POST | /payment/create_user | 1551 | 0 | 60 | 160 | 350 | 82 | 16 | 542 | 51 | 3.9 | 0 |
| POST | /stock/add{item_id}/{number} | 2254 | 0 | 760 | 1800 | 2400 | 826 | 18 | 2877 | 0 | 6.9 | 0 |
| GET | /stock/find{item_id} | 133 | 0 | 760 | 1700 | 2600 | 821 | 9 | 2741 | 39 | 0.5 | 0 |
| POST | /stock/item/create{price} | 2283 | 0 | 760 | 1800 | 2400 | 836 | 16 | 2907 | 51 | 7.5 | 0 |
| POST | /stock/subtract{item_id}/{number} | 133 | 0 | 720 | 1800 | 2400 | 828 | 23 | 2443 | 0 | 0.5 | 0 |
| Aggregated | | 13343 | 0 | 310 | 1700 | 3100 | 637 | 9 | 6728 | 24 | 42.6 | 0 |

Q&A